

Jay Prakash

Chennai, India • Open to Relocation • jaysingha261@gmail.com • +917033762663 • [LinkedIn](#) • [GitHub](#) • [LeetCode](#)

Summary

Backend Engineer with 2 years building and operating Java Spring Boot microservices across a 17-service production healthtech platform on GCP. Delivered performance improvements from minutes to milliseconds through query precomputation and Redis caching. Additional hands-on work in LangGraph and RAG pipelines.

Skills

Languages: Java, Python, SQL, JavaScript

Backend: Spring Boot, Microservices, REST APIs, Hibernate/JPA, Keycloak, FastAPI

Caching: Redis

Databases: PostgreSQL (pgvector), ClickHouse, BigQuery

Cloud & DevOps: GCP, GKE, Kubernetes, Docker, Helm, Jenkins, GitHub Actions

AI/ML: LangGraph, LangChain, RAG, LLM APIs (Groq, Gemini), LLM-as-a-Judge, Langfuse

Experience

Mindsprint

Software Engineer

Jun 2024 – Present

- Promoted from Junior Engineer to Software Engineer within 18 months in recognition of technical contributions and ownership.
- Developed and maintained Java Spring Boot microservices across a platform comprising 17 services, enabling core business workflows, third-party integrations, and data processing for healthcare clients.
- Built and deployed a scheduled CronJob to precompute retailer-wise patient metrics, eliminating repeated database-intensive calculations and improving response times from 2–3 minutes to milliseconds.
- Added Redis caching for frequently accessed data, cutting API response time from 3s to 500ms and lowering database calls by 20–30%.

Projects

Code Review Agent | [Demo](#) | [GitHub](#)

- Built a GitHub App that reviews pushed commits automatically: webhook receiver → diff extraction → LLM review via Groq → inline commit comment.
- Deployed as a FastAPI service on Render with uptime monitoring and automated health checks.

Agentic AI RAG System | [Live](#) | [GitHub](#)

- Built an agentic RAG pipeline in LangGraph where a multi-tool router dispatches queries to hybrid vector/full-text retrieval, web search, or the GitHub API.
- Added an LLM-as-judge loop that detects ungrounded answers, auto-rewrites the query up to 3 times, then escalates to a human review queue that feeds approved answers back into the knowledge base.
- Shipped as a FastAPI service with Langfuse tracing, gated by a 113-test mocked-LLM pytest suite in CI, deployed on Hugging Face Spaces with Neon Postgres.

Education

SRM University

B.Tech, Computer Science, CGPA: 9.25, Mar 2024